

Código, implementación y discusión del problema de reconocimiento de dígitos

Integrantes del grupo:

- Juan David Bernal Vesga
- Oscar David Ordóñez Bolaños
- Juan Pablo Urrutia Parrado

A continuación, se muestra el paso a paso de la implementación de la solución, así como también la discusión de algunos resultados

Para ejecutar el cuaderno, se deben cumplir los siguientes requerimientos:

- Para este propósito se usa `Python 3.10.0` como la versión de ejecución.
- Se deben tener las librerías instaladas, las cuales se encuentran en `requirements.txt`

Nota: para instalar las librerías, ejecutar en la terminal el siguiente comando: `pip install -r requirements.txt`

Dataset utilizado

Primeramente, se importan las librerías usadas para nuestro propósito, y se definen algunas funciones que sirven como herramientas para el análisis.

```
In [ ]: from keras.datasets import mnist
import matplotlib.pyplot as plt
import keras
import numpy as np
```

El dataset que se usa es el `mnist`, el cual es un dataset de 60.000 entradas para entrenar y 10.000 para validar, como se muestra a continuación:

```
In [ ]: # Import dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# Examine dataset shapes
print(x_train.shape, x_test.shape)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 1s 0us/step
(60000, 28, 28) (10000, 28, 28)
```


Visualizando algunos datos del dataset, vemos que son de tamaño 28x28 y están centrados, como se muestra a continuación:

```
In [ ]: fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2)
fig.suptitle('Dataset examples')
```

```

test = np.random.randint(0,10000)
ax1.imshow(x_test[test], cmap="binary")
ax1.title.set_text(y_test[test])
ax1.axis('off')
ax2.imshow(x_test[test+1], cmap="binary")
ax2.title.set_text(y_test[test+1])
ax2.axis('off')
ax3.imshow(x_test[test+2], cmap="binary")
ax3.title.set_text(y_test[test+2])
ax3.axis('off')
ax4.imshow(x_test[test+3], cmap="binary")
ax4.title.set_text(y_test[test+3])
ax4.axis('off');

```



Ahora que se ha visto la estructura de los datos, los preparamos para ser usados por el modelo.

Creación del modelo inicial

Para trabajar con la API de Keras, se debe trabajar con imágenes de formato (M x N x 1). Se usa el método `.reshape()` para realizar esta acción. Finalmente, se normalizan los datos de la imagen dividiendo cada valor de píxel por 255 (ya que los valores RGB pueden variar de 0 a 255). Además, se necesita convertir la variable de verificación de `int` a categorías, ya que se trabajan con categorías del 0 al 9. Esto se puede lograr con la función `to_categorical()`:

```

In [ ]: # Reshape
x_train = x_train.reshape(x_train.shape[0], 28, 28, 1)
x_test = x_test.reshape(x_test.shape[0], 28, 28, 1)

# Parse and normalize
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255

# Turn Y into categorical data
y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)

```



```
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

model.save("src/models/og_model.h5")
```

Epoch 1/20
469/469 [=====] - 18s 10ms/step - loss: 0.2418 - accuracy: 0.9261 - val_loss: 0.0511 - val_accuracy: 0.9839
Epoch 2/20
469/469 [=====] - 4s 9ms/step - loss: 0.0836 - accuracy: 0.9747 - val_loss: 0.0395 - val_accuracy: 0.9876
Epoch 3/20
469/469 [=====] - 4s 9ms/step - loss: 0.0643 - accuracy: 0.9809 - val_loss: 0.0312 - val_accuracy: 0.9901
Epoch 4/20
469/469 [=====] - 4s 9ms/step - loss: 0.0512 - accuracy: 0.9846 - val_loss: 0.0293 - val_accuracy: 0.9898
Epoch 5/20
469/469 [=====] - 5s 10ms/step - loss: 0.0442 - accuracy: 0.9863 - val_loss: 0.0283 - val_accuracy: 0.9909
Epoch 6/20
469/469 [=====] - 4s 9ms/step - loss: 0.0376 - accuracy: 0.9882 - val_loss: 0.0258 - val_accuracy: 0.9914
Epoch 7/20
469/469 [=====] - 5s 10ms/step - loss: 0.0346 - accuracy: 0.9888 - val_loss: 0.0264 - val_accuracy: 0.9912
Epoch 8/20
469/469 [=====] - 5s 10ms/step - loss: 0.0298 - accuracy: 0.9906 - val_loss: 0.0269 - val_accuracy: 0.9907
Epoch 9/20
469/469 [=====] - 4s 9ms/step - loss: 0.0290 - accuracy: 0.9907 - val_loss: 0.0268 - val_accuracy: 0.9919
Epoch 10/20
469/469 [=====] - 4s 9ms/step - loss: 0.0254 - accuracy: 0.9917 - val_loss: 0.0253 - val_accuracy: 0.9926
Epoch 11/20
469/469 [=====] - 4s 9ms/step - loss: 0.0225 - accuracy: 0.9925 - val_loss: 0.0265 - val_accuracy: 0.9922
Epoch 12/20
469/469 [=====] - 5s 10ms/step - loss: 0.0220 - accuracy: 0.9930 - val_loss: 0.0257 - val_accuracy: 0.9922
Epoch 13/20
469/469 [=====] - 4s 9ms/step - loss: 0.0198 - accuracy: 0.9933 - val_loss: 0.0263 - val_accuracy: 0.9915
Epoch 14/20
469/469 [=====] - 4s 9ms/step - loss: 0.0202 - accuracy: 0.9936 - val_loss: 0.0261 - val_accuracy: 0.9925
Epoch 15/20
469/469 [=====] - 5s 10ms/step - loss: 0.0191 - accuracy: 0.9936 - val_loss: 0.0307 - val_accuracy: 0.9912
Epoch 16/20
469/469 [=====] - 5s 12ms/step - loss: 0.0171 - accuracy: 0.9947 - val_loss: 0.0275 - val_accuracy: 0.9922
Epoch 17/20
469/469 [=====] - 5s 10ms/step - loss: 0.0163 - accuracy: 0.9947 - val_loss: 0.0295 - val_accuracy: 0.9922
Epoch 18/20
469/469 [=====] - 4s 9ms/step - loss: 0.0146 - accuracy: 0.9952 - val_loss: 0.0283 - val_accuracy: 0.9929
Epoch 19/20
469/469 [=====] - 4s 9ms/step - loss: 0.0160 - accuracy: 0.9946 - val_loss: 0.0326 - val_accuracy: 0.9925
Epoch 20/20
469/469 [=====] - 4s 9ms/step - loss: 0.0137 - accuracy: 0.9954 - val_loss: 0.0293 - val_accuracy: 0.9931


Test loss: 0.029333407059311867
Test accuracy: 0.9930999875068665

```
In [ ]: original_model = keras.models.load_model('src/models/og_model.h5')
```

A continuación, una prueba sobre un dato aislado:

```
In [ ]: import imageio

# Import and process the image
im = imageio.imread("img/2.png")
plt.imshow(im, cmap = plt.get_cmap('gray'))
plt.show()
gray = np.dot(im[..., :3], [0.299, 0.587, 0.114])
gray = gray.reshape(1, 28, 28, 1)
gray /= 255
```



```
In [ ]: # Use the model to predict
prediction = model.predict(gray)
print(prediction.argmax())
```

```
1/1 [=====] - 0s 17ms/step
2
```

Como vemos, el modelo predice correctamente sobre la imagen suministrada.

Problemas de Overfitting

Haciendo un análisis sobre el porcentaje de error del modelo anterior, se puede ver lo siguiente:

```
In [ ]: # Aux. function to graph results
def eval_metric(model, history, metric_name, name):
    """
    Function to evaluate a trained model on a chosen metric.
    Training and validation metric are plotted in a
    line chart for each epoch.

    Parameters:
        history : model training history
        metric_name : loss or accuracy


    Output:
        line chart with epochs of x-axis and metric on
```

```

    ... y-axis
    ...
    metric = history.history[metric_name]
    val_metric = history.history['val_' + metric_name]
    e = range(1, 20 + 1)
    plt.plot(e, metric, label='Entrenamiento ' + metric_name)
    plt.plot(e, val_metric, label='Validación ' + metric_name)
    plt.xlabel('Epoch')
    plt.ylabel(metric_name)
    plt.title('Comparación ' + metric_name + ' de entrenamiento y validación par
    plt.legend()
    plt.show()

```

```
In [ ]: eval_metric(original_model, og_model_hist, 'loss', 'Original')
```



El modelo original presenta síntomas de 'overfitting'. Esto es, cuando el modelo queda sobre-entrenado sobre un dataset, tanto que es incapaz de reconocer dígitos de manera generalizada. Esto se puede detectar encontrando que a medida que se entrena el modelo, el loss (porcentaje de error) sobre el dataset de entreno baja, mientras que el loss sobre el dataset de validación se mantiene o sube. Para la primera implementación del modelo, se implementan 'Dropout layers'; sin embargo, parece no ser suficiente para resolver el problema. Aunque no es muy grave, se puede ver como el loss sobre el set de validación deja de disminuir en el octavo entrenamiento, y luego sube. Para atacar este problema, se plantean algunas estrategias:

1. Reducir el número de neuronas de la capa de segunda convolucion de 64 a 32, y de la capa oculta de 128 a 64.

Así se reduce la capacidad de la red y se obliga a aprender los patrones que importan o que minimizan la pérdida.

```

In [ ]: # Create layers as explained before
model_1 = Sequential()
model_1.add(Conv2D(32, kernel_size=(3, 3),
    activation='relu',
    input_shape=(28, 28, 1)))

model_1.add(Conv2D(32, (3, 3), activation='relu')) ## 64 -> 32
model_1.add(MaxPooling2D(pool_size=(2, 2)))

```

```

model_1.add(Dropout(0.25))

model_1.add(Flatten())

model_1.add(Dense(64, activation='relu'))           ## 128 ->64
model_1.add(Dropout(0.5))
model_1.add(Dense(10, activation='softmax'))

# Compile the model_1 using adam as optimizer and categorical_crossentropy as th
model_1.compile(loss=keras.losses.categorical_crossentropy,
                 optimizer='adam',
                 metrics=['accuracy'])

```

```

In [ ]: batch_size = 128
        epochs = 20

        small_model_hist = model_1.fit(x_train, y_train,
                                       batch_size=batch_size,
                                       epochs=epochs,
                                       verbose=1,
                                       validation_data=(x_test, y_test))
        score = model_1.evaluate(x_test, y_test, verbose=0)
        print('Test loss:', score[0])
        print('Test accuracy:', score[1])

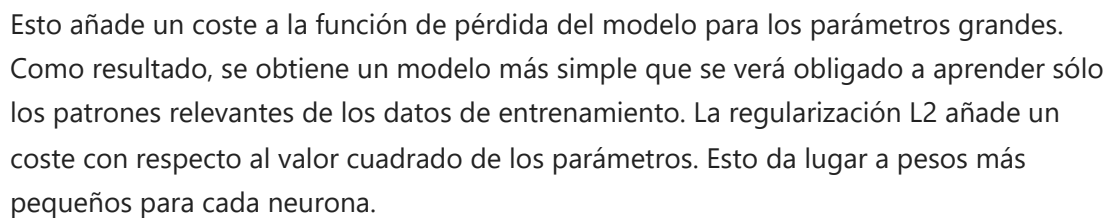
        model_1.save("src/models/small_model.h5")

        eval_metric(model_1, small_model_hist, 'loss', 'Original')

```


Epoch 1/20
469/469 [=====] - 4s 8ms/step - loss: 0.3600 - accuracy: 0.8884 - val_loss: 0.0731 - val_accuracy: 0.9745
Epoch 2/20
469/469 [=====] - 3s 7ms/step - loss: 0.1416 - accuracy: 0.9584 - val_loss: 0.0464 - val_accuracy: 0.9854
Epoch 3/20
469/469 [=====] - 3s 7ms/step - loss: 0.1083 - accuracy: 0.9680 - val_loss: 0.0388 - val_accuracy: 0.9867
Epoch 4/20
469/469 [=====] - 3s 7ms/step - loss: 0.0911 - accuracy: 0.9725 - val_loss: 0.0361 - val_accuracy: 0.9885
Epoch 5/20
469/469 [=====] - 3s 7ms/step - loss: 0.0816 - accuracy: 0.9760 - val_loss: 0.0329 - val_accuracy: 0.9894
Epoch 6/20
469/469 [=====] - 3s 7ms/step - loss: 0.0735 - accuracy: 0.9779 - val_loss: 0.0320 - val_accuracy: 0.9893
Epoch 7/20
469/469 [=====] - 3s 7ms/step - loss: 0.0662 - accuracy: 0.9795 - val_loss: 0.0300 - val_accuracy: 0.9897
Epoch 8/20
469/469 [=====] - 3s 7ms/step - loss: 0.0614 - accuracy: 0.9807 - val_loss: 0.0296 - val_accuracy: 0.9899
Epoch 9/20
469/469 [=====] - 3s 7ms/step - loss: 0.0569 - accuracy: 0.9825 - val_loss: 0.0283 - val_accuracy: 0.9904
Epoch 10/20
469/469 [=====] - 3s 7ms/step - loss: 0.0525 - accuracy: 0.9838 - val_loss: 0.0287 - val_accuracy: 0.9910
Epoch 11/20
469/469 [=====] - 3s 7ms/step - loss: 0.0514 - accuracy: 0.9843 - val_loss: 0.0286 - val_accuracy: 0.9916
Epoch 12/20
469/469 [=====] - 3s 7ms/step - loss: 0.0487 - accuracy: 0.9849 - val_loss: 0.0336 - val_accuracy: 0.9897
Epoch 13/20
469/469 [=====] - 3s 7ms/step - loss: 0.0460 - accuracy: 0.9851 - val_loss: 0.0287 - val_accuracy: 0.9914
Epoch 14/20
469/469 [=====] - 3s 7ms/step - loss: 0.0445 - accuracy: 0.9852 - val_loss: 0.0290 - val_accuracy: 0.9916
Epoch 15/20
469/469 [=====] - 3s 7ms/step - loss: 0.0416 - accuracy: 0.9866 - val_loss: 0.0319 - val_accuracy: 0.9907
Epoch 16/20
469/469 [=====] - 3s 7ms/step - loss: 0.0397 - accuracy: 0.9871 - val_loss: 0.0287 - val_accuracy: 0.9925
Epoch 17/20
469/469 [=====] - 3s 7ms/step - loss: 0.0388 - accuracy: 0.9875 - val_loss: 0.0298 - val_accuracy: 0.9922
Epoch 18/20
469/469 [=====] - 3s 7ms/step - loss: 0.0353 - accuracy: 0.9880 - val_loss: 0.0259 - val_accuracy: 0.9918
Epoch 19/20
469/469 [=====] - 3s 7ms/step - loss: 0.0341 - accuracy: 0.9884 - val_loss: 0.0263 - val_accuracy: 0.9925
Epoch 20/20
469/469 [=====] - 3s 7ms/step - loss: 0.0362 - accuracy: 0.9881 - val_loss: 0.0317 - val_accuracy: 0.9903

Test accuracy: 0.9902999997138977

[illegible]

```
score = model_2.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])


model_2.save("src/models/conv_reg_model.h5")

eval_metric(model_2, reg_model_hist, 'loss', 'Original')
```

Epoch 1/20
469/469 [=====] - 5s 10ms/step - loss: 0.2832 - accuracy: 0.9190 - val_loss: 0.0735 - val_accuracy: 0.9829
Epoch 2/20
469/469 [=====] - 5s 10ms/step - loss: 0.1113 - accuracy: 0.9717 - val_loss: 0.0521 - val_accuracy: 0.9879
Epoch 3/20
469/469 [=====] - 4s 10ms/step - loss: 0.0833 - accuracy: 0.9779 - val_loss: 0.0481 - val_accuracy: 0.9881
Epoch 4/20
469/469 [=====] - 4s 10ms/step - loss: 0.0702 - accuracy: 0.9811 - val_loss: 0.0403 - val_accuracy: 0.9900
Epoch 5/20
469/469 [=====] - 6s 12ms/step - loss: 0.0614 - accuracy: 0.9838 - val_loss: 0.0366 - val_accuracy: 0.9897
Epoch 6/20
469/469 [=====] - 7s 14ms/step - loss: 0.0563 - accuracy: 0.9847 - val_loss: 0.0381 - val_accuracy: 0.9893
Epoch 7/20
469/469 [=====] - 7s 14ms/step - loss: 0.0508 - accuracy: 0.9862 - val_loss: 0.0356 - val_accuracy: 0.9903
Epoch 8/20
469/469 [=====] - 6s 13ms/step - loss: 0.0436 - accuracy: 0.9875 - val_loss: 0.0343 - val_accuracy: 0.9908
Epoch 9/20
469/469 [=====] - 4s 9ms/step - loss: 0.0405 - accuracy: 0.9893 - val_loss: 0.0344 - val_accuracy: 0.9905
Epoch 10/20
469/469 [=====] - 4s 9ms/step - loss: 0.0374 - accuracy: 0.9896 - val_loss: 0.0331 - val_accuracy: 0.9921
Epoch 11/20
469/469 [=====] - 4s 9ms/step - loss: 0.0345 - accuracy: 0.9904 - val_loss: 0.0320 - val_accuracy: 0.9923
Epoch 12/20
469/469 [=====] - 5s 10ms/step - loss: 0.0349 - accuracy: 0.9899 - val_loss: 0.0346 - val_accuracy: 0.9904
Epoch 13/20
469/469 [=====] - 4s 9ms/step - loss: 0.0311 - accuracy: 0.9912 - val_loss: 0.0312 - val_accuracy: 0.9925
Epoch 14/20
469/469 [=====] - 4s 9ms/step - loss: 0.0299 - accuracy: 0.9914 - val_loss: 0.0330 - val_accuracy: 0.9916
Epoch 15/20
469/469 [=====] - 4s 9ms/step - loss: 0.0273 - accuracy: 0.9925 - val_loss: 0.0332 - val_accuracy: 0.9919
Epoch 16/20
469/469 [=====] - 4s 9ms/step - loss: 0.0261 - accuracy: 0.9928 - val_loss: 0.0328 - val_accuracy: 0.9912
Epoch 17/20
469/469 [=====] - 5s 10ms/step - loss: 0.0241 - accuracy: 0.9930 - val_loss: 0.0347 - val_accuracy: 0.9926
Epoch 18/20
469/469 [=====] - 4s 9ms/step - loss: 0.0248 - accuracy: 0.9931 - val_loss: 0.0330 - val_accuracy: 0.9914
Epoch 19/20
469/469 [=====] - 5s 10ms/step - loss: 0.0232 - accuracy: 0.9934 - val_loss: 0.0317 - val_accuracy: 0.9929
Epoch 20/20
469/469 [=====] - 6s 12ms/step - loss: 0.0215 - accuracy: 0.9938 - val_loss: 0.0390 - val_accuracy: 0.9914

Test loss: 0.03898705914616585

Test accuracy: 0.9914000034332275



Al combinar las anteriores dos soluciones, se obtiene la solución final:

```
In [ ]: # Create layers as explained before
model_3 = Sequential()
model_3.add(Conv2D(32, kernel_size=(3, 3),
                  activation='relu',
                  input_shape=(28, 28, 1), kernel_regularizer=keras.regularizers.l2(0.01), bias_regularizer=keras.regularizers.l2(0.01)))

model_3.add(Conv2D(32, (3, 3), activation='relu'))          ## 64 -> 32
model_3.add(MaxPooling2D(pool_size=(2, 2)))

model_3.add(Dropout(0.25))

model_3.add(Flatten())

model_3.add(Dense(64, activation='relu'))                  ## 128 -> 64
model_3.add(Dropout(0.5))
model_3.add(Dense(10, activation='softmax'))

# Compile the model_3 using adam as optimizer and categorical_crossentropy as the loss function
model_3.compile(loss=keras.losses.categorical_crossentropy,
                optimizer='adam',
                metrics=['accuracy'])
```

```
In [ ]: batch_size = 128
epochs = 20

small_conv_reg_hist = model_3.fit(x_train, y_train,
                                batch_size=batch_size,
                                epochs=epochs,
                                verbose=1,
                                validation_data=(x_test, y_test))
score = model_3.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])


model_3.save("src/models/small_conv_reg_model.h5")

eval_metric(model_3, small_conv_reg_hist, 'loss', 'Original')
```

Epoch 1/20
469/469 [=====] - 4s 8ms/step - loss: 0.4071 - accuracy: 0.8810 - val_loss: 0.1034 - val_accuracy: 0.9716
Epoch 2/20
469/469 [=====] - 3s 7ms/step - loss: 0.1783 - accuracy: 0.9521 - val_loss: 0.0714 - val_accuracy: 0.9822
Epoch 3/20
469/469 [=====] - 4s 8ms/step - loss: 0.1402 - accuracy: 0.9621 - val_loss: 0.0569 - val_accuracy: 0.9849
Epoch 4/20
469/469 [=====] - 4s 8ms/step - loss: 0.1206 - accuracy: 0.9678 - val_loss: 0.0537 - val_accuracy: 0.9857
Epoch 5/20
469/469 [=====] - 3s 7ms/step - loss: 0.1065 - accuracy: 0.9711 - val_loss: 0.0474 - val_accuracy: 0.9875
Epoch 6/20
469/469 [=====] - 3s 7ms/step - loss: 0.0964 - accuracy: 0.9744 - val_loss: 0.0435 - val_accuracy: 0.9886
Epoch 7/20
469/469 [=====] - 3s 7ms/step - loss: 0.0896 - accuracy: 0.9754 - val_loss: 0.0432 - val_accuracy: 0.9888
Epoch 8/20
469/469 [=====] - 3s 7ms/step - loss: 0.0835 - accuracy: 0.9777 - val_loss: 0.0439 - val_accuracy: 0.9887
Epoch 9/20
469/469 [=====] - 3s 7ms/step - loss: 0.0821 - accuracy: 0.9770 - val_loss: 0.0417 - val_accuracy: 0.9890
Epoch 10/20
469/469 [=====] - 4s 9ms/step - loss: 0.0746 - accuracy: 0.9786 - val_loss: 0.0433 - val_accuracy: 0.9883
Epoch 11/20
469/469 [=====] - 3s 7ms/step - loss: 0.0729 - accuracy: 0.9794 - val_loss: 0.0394 - val_accuracy: 0.9900
Epoch 12/20
469/469 [=====] - 3s 7ms/step - loss: 0.0687 - accuracy: 0.9810 - val_loss: 0.0451 - val_accuracy: 0.9894
Epoch 13/20
469/469 [=====] - 3s 7ms/step - loss: 0.0679 - accuracy: 0.9810 - val_loss: 0.0367 - val_accuracy: 0.9899
Epoch 14/20
469/469 [=====] - 3s 7ms/step - loss: 0.0649 - accuracy: 0.9818 - val_loss: 0.0369 - val_accuracy: 0.9896
Epoch 15/20
469/469 [=====] - 3s 7ms/step - loss: 0.0622 - accuracy: 0.9823 - val_loss: 0.0361 - val_accuracy: 0.9900
Epoch 16/20
469/469 [=====] - 3s 7ms/step - loss: 0.0598 - accuracy: 0.9830 - val_loss: 0.0355 - val_accuracy: 0.9909
Epoch 17/20
469/469 [=====] - 3s 7ms/step - loss: 0.0579 - accuracy: 0.9830 - val_loss: 0.0366 - val_accuracy: 0.9899
Epoch 18/20
469/469 [=====] - 3s 7ms/step - loss: 0.0550 - accuracy: 0.9839 - val_loss: 0.0341 - val_accuracy: 0.9914
Epoch 19/20
469/469 [=====] - 3s 7ms/step - loss: 0.0543 - accuracy: 0.9845 - val_loss: 0.0352 - val_accuracy: 0.9904
Epoch 20/20
469/469 [=====] - 3s 7ms/step - loss: 0.0554 - accuracy: 0.9842 - val_loss: 0.0329 - val_accuracy: 0.9914

Test loss: 0.032890982925891876

Test accuracy: 0.9914000034332275



Vemos entonces que el modelo no presenta 'overfitting' y está listo para ser usado por el usuario final.